

Stochastic gradient descent for SVM

Lluís Garrido – lluis.garrido@ub.edu

December 2021

Abstract

This laboratory is focused on stochastic subgradient methods in the support vector machines. These methods have been introduced in the Deep Learning course and here we will see how they are applied to train a Support Vector Machine.

1 Support Vector Machines

In the previous lab we have formulated pattern classification using linear functions to provide the characterization of a sample. We recall here the main aspects of the formulation just as a reminder. Suppose we have a set of m training data, $x_i \in R^n$, with classification y_i , where either $y_i = 1$ or $y_i = -1$ (the data point has a certain property or not), see Figure 1 on the left. Suppose it is possible to find some hyperplane $\mathbf{w}^T \mathbf{x} + b = 0$, $\mathbf{w} \in R^n$ and $b \in R$, which separates the positive points from the negative. The problem is to determine the coefficients \mathbf{w} and b that solve

$$\begin{aligned} \text{mimize} \quad & f(\mathbf{w}, b) = \frac{1}{2} \mathbf{w}^T \mathbf{w} \\ \text{subject to} \quad & y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 \quad i = 1 \dots m \end{aligned}$$

Once the coefficients \mathbf{w} and b of the separating hyperplane are found from the training data, we can use the value of the function $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$ (our “learning machine”) to predict whether a new point $\bar{\mathbf{x}}$ has the property of interest or not, depending on the sign of $f(\bar{\mathbf{x}})$. The previous formulation assumes that the data is separable, that is, there exists a hyperplane separating the positive points from the negative points exists. In the previous lab, we presented a formulation that allowed to transform the problem to the dual space

For the case where the data set is not separable, we can refine the approach of the separable case, see Figure 1. We will now allow the points to violate the equations of the separating hyperplane, but we will impose a penalty for the violation. Letting the nonnegative variable ξ_i denote the amount by which the point x_i violates the constraint at the margin, we now add to the objective a term proportional to the sum of the violations. The added penalty term takes the form $K \sum \xi_i$ and is added to the objective, where the larger the value of the parameter K , the larger the penalty for violating the separation. Our problem is now to find \mathbf{w} , b and ξ that solve

$$\begin{aligned} \text{mimize} \quad & f(\mathbf{w}, b, \xi) = \frac{1}{2} \mathbf{w}^T \mathbf{w} + K \sum \xi_i \\ \text{subject to} \quad & y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i \quad i = 1 \dots m \\ & \xi_i \geq 0 \end{aligned} \tag{1}$$

This is the original or *primal* problem. In the previous lab we have seen that many optimization problems have a companion problem called the dual problem. The dual problem may be easier to




Figure 1: Linear separating hyperplane for the non separable case. Image taken from Griva, I.; Nash, S.; Sofer, A., “Linear and nonlinear optimization”, SIAM.

solve, and if the optimal solution to the dual problem is known, then (in nondegenerate cases) the optimal solution to the primal problem can be easily computed. In addition, for the case of the Support Vector Machine, the dual optimization problem can be written in terms of dot products, thereby making it possible to use kernel functions.

These reasons are not a limitation for solving the problem in the primal, mainly by writing the optimization problem as an unconstrained one and by using the Representer theorem to be able to deal with the non-linear case. Intuitively, the primal optimization should be superior than the dual because it directly minimizes the quantity we are interested in, \mathbf{w} , whereas in the dual problem we solve the dual variable, α (see previous lab).

The primal problem of Equation (1) can be rewritten in an unconstrained way as

$$f(\mathbf{w}, b) = \frac{1}{2} \mathbf{w}^T \mathbf{w} + K \sum_{i=0}^m \max(0, 1 - y_i(\mathbf{w}^T \mathbf{x}_i + b)) \quad (2)$$

This is the function to be minimized and it can be solved using, for instance, the gradient or the Newton method. We now rewrite a bit the previous equation in order to avoid too “large numbers” when doing numerical computations. The previous equation is rewritten as

$$f(\mathbf{w}, b) = \frac{\lambda}{2} \mathbf{w}^T \mathbf{w} + \sum_{i=0}^m \max(0, 1 - y_i(\mathbf{w}^T \mathbf{x}_i + b))$$

where λ should be a “small value” (since usually K is a large value).

In this case we propose to use the Stochastic gradient algorithm. The Stochastic gradient descent is a drastic simplification of the gradient descent. Instead of computing the gradient of Equation (2) using all the terms of the sum, $i = 1 \dots m$, each iteration estimates the gradient of $f(\mathbf{w}, b)$ on the basis of a single randomly picked example \mathbf{x}_t , where t is the stochastic process $t = 1, \dots$. The stochastic gradient algorithm is

$$\begin{aligned}\mathbf{w}_{t+1} &\leftarrow \mathbf{w}_t - \gamma_t \left(\lambda \mathbf{w}_t + \begin{cases} 0 & \text{if } y_t(\mathbf{w}_t^T \mathbf{x}_t + b) > 1 \\ -y_t \mathbf{x}_t & \text{otherwise} \end{cases} \right) \\ b_{t+1} &\leftarrow b_t - \gamma_t \begin{cases} 0 & \text{if } y_t(\mathbf{w}_t^T \mathbf{x}_t + b) > 1 \\ -y_t & \text{otherwise} \end{cases}\end{aligned}$$

Observe that the values of \mathbf{w}_t and b_t are updated even if $f(\mathbf{w}_{t+1}, b_{t+1})$ increases with respect to $f(\mathbf{w}_t, b_t)$. The convergence of the stochastic gradient descent has been studied extensively in the literature. Convergence results usually require decreasing gains γ_t at each iteration. The convergence speed of stochastic gradient descent is in fact limited by the noisy approximation of the true gradient. Under sufficient regularity conditions, the best convergence speed is achieved using gains $\gamma_t \sim t^{-1}$.

For the case of the mini-batch a similar procedure is followed: assume you want to perform a mini-batch of a set S of M elements. At each iteration t you select M elements from the m available samples. The stochastic minibatch algorithm is

$$\begin{aligned}\mathbf{w}_{t+1} &\leftarrow \mathbf{w}_t - \gamma_t \left(\lambda \mathbf{w} + \sum_{k \in S} \begin{cases} 0 & \text{if } y_k(\mathbf{w}^T \mathbf{x}_k + b) > 1 \\ -y_k \mathbf{x}_k & \text{otherwise} \end{cases} \right) \\ b_{t+1} &\leftarrow b_t - \gamma_t \left(\sum_{k \in S} \begin{cases} 0 & \text{if } y_k(\mathbf{w}^T \mathbf{x}_k + b) > 1 \\ -y_k & \text{otherwise} \end{cases} \right)\end{aligned}$$

Proposed experiments

For the following experiments, it is interesting to plot the classification line at level lines 0, +1 and -1. If you want to compare with the dual method of the previous lab, make sure you take the same value of K (in this case λ) in both methods.

You are recommended to use separable sets to perform such experiments. However, the algorithm should also work for the non-separable case.

1. Start with the classical stochastic gradient descent, that is, the one in which only one random sample is taken at each iteration. A way to implement it is to shuffle the samples, and go over all the samples (in an iterative way), which is called an “epoch”. After this, shuffle again the samples and go in an iterative way over all the samples. You are recommended to start with a relatively small value of γ_t , for instance, $\gamma_t = 1/\mu$, with μ initialized to $\mu = 100$, for instance. The value of μ is increased after each iteration. You may increase μ after each iteration by a value of 1, 0.5 or 0.1, for instance. You also may use $\lambda = 10^{-4}$ or $\lambda = 10^{-2}$ for instance. Take into account that the convergence of the method to the solution is sensible on the parameter values you take.
2. Perform some experiments with the mini-batch. You may try different number of samples to construct the mini-batches, e.g. 10 or 50 samples. In order to simplify the problem and be able to perform comparisons between the stochastic and minibatch, take a number of samples of 200 (100 at each class, as performed in the previous lab) and mini-batches with a size that of 10, 20 or 50, for instance. After each epoch you perform a shuffle of the samples and construct again the mini-batches.

3. Perform a plot of the logarithm of $f(\mathbf{w}, b)$ along each epoch (it is important to plot the logarithm of the function) for the stochastic gradient descent, minibatch and gradient descent (which is equivalent to use a mini-batch with a size equal to the number of samples). This allows you to see how fast each of the methods approaches the optimal solution and perform a comparison between the different experiments you have done.

Report

The report can be done in *pairs of two persons* (or *individually*). You are asked to deliver a only one report (PDF, notebook, or whatever else you prefer) of the work you have performed in this and previous lab. You also are requested to perform a comparison between the dual method studied in the previous lab, as well as the stochastic and the mini batch using the primal equation (2). Take into account that the problem you are solving is quadratic, convex, and thus it (mathematically) only has only minimum. However, the numerical method you use may have difficulties to arrive to such minimum since the convergence of the method to the solution is sensible on the parameter values you take.

Notes

There is a great paper that treats the problem of training a Support Vector Machine in the primal. It treats the problem of training using Newton optimization. It is “Training a Support Vector Machine in the Primal”, by O. Chapelle, 2006. For the stochastic method, you may take a look at “Large-Scale Machine Learning with Stochastic Gradient Descent”, by L. Bottou, 2010.